Qwen3.5:4B prompting master guide: definitive engineering and integration for local SLM deployment

Processing Iteration 1/10: Validating core architectural parameters and release context. The emergence of Alibaba's Qwen3.5:4B on March 3, 2026, represents a transformative moment in the Small Language Model (SLM) ecosystem, achieving what early observers describe as an unprecedented "intelligence density".[1] By integrating a sophisticated hybrid attention mechanism and native multimodal processing into a 4-billion parameter footprint, Alibaba has established a new benchmark for on-device AI capabilities.[3] This model is not merely a shrunken version of its larger 397B flagship sibling; it is an architecturally optimized system designed specifically to overcome the structural limits that historically hindered SLMs, such as quadratic attention scaling and the performance degradation associated with adapter-based vision encoders.[5] As the mid-tier entry in the Qwen3.5 Small series, the 4B variant serves as a stable foundational model for lightweight agents, capable of navigating graphical user interfaces, interpreting dense documentation, and maintaining logical consistency over long-horizon reasoning tasks that previously required models eight to ten times its size.[2]

Executive summary: Qwen3.5 4B architecture and best use cases

Processing Iteration 2/10: Analyzing the hybrid attention and MoE layout for performance optimization. The technical foundation of Qwen3.5:4B is built upon a hybrid architecture that combines Gated Delta Networks (DeltaNet) with sparse Mixture-of-Experts (MoE) and traditional softmax attention.[5] A critical innovation is the 3:1 ratio of linear attention blocks to full softmax attention blocks, a design choice that addresses the memory and compute bottlenecks of the standard Transformer architecture.[8] This implementation allows the model to selectively retain or forget information within a bounded hidden state, which facilitates a native context window of 262,144 tokens that can be extended up to 1,010,000 tokens through RoPE scaling techniques such as YaRN.[6] Such an extensive context window is a significant departure for models in the 4B class, positioning Qwen3.5:4B as a viable tool for long-document analysis, complex code review, and multi-turn agentic interactions where context retention is paramount.[11]

The MoE component of the 4B model activates only a fraction of its total parameters during each forward pass, which maintains inference efficiency and reduces the computational load on consumer-grade hardware.[5] For example, while the model has 4 billion dense parameters, the sparse activation ensures that throughput remains high even on edge devices such as

high-end smartphones or local laptops equipped with M4 or RTX 40-series GPUs.[14] This architectural efficiency is complemented by a vocabulary of 248,320 tokens, trained specifically to enhance encoding for over 201 languages and dialects.[6] The combination of these features allows the 4B model to match or exceed the performance of older 30B to 72B parameter models in reasoning and tool-invocation benchmarks.[2]

Architectural specifications and hardware requirements

Optimal deployment of Qwen3.5:4B requires a nuanced understanding of its memory footprint across various quantization levels. Community benchmarks from March 2026 indicate that the model is highly resilient to 4-bit and 6-bit quantization, with minimal loss in reasoning accuracy.[8]

Hardware Tier	Quantization	Memory Required (VRAM/RAM)	Inference Speed (Est.)
Mobile Edge (iPhone 15+)	Q4_0 / GGUF	2.6 GB - 3.5 GB [14]	8-12 tokens/sec
Standard Desktop PC	Q4_K_M / GGUF	4.5 GB - 5.5 GB [17]	25-40 tokens/sec
Professional Workstation	Q8_0 / BF16	7.9 GB - 14 GB [8]	60+ tokens/sec
GPU Servers (H100/A100)	FP8 / BF16	14 GB+ [17]	300+ tokens/sec

Beyond memory, the model's DeltaNet implementation necessitates a minimum input of 2 tokens to function correctly due to its Conv1d kernel size requirement; a single-token input will trigger a crash in several inference backends, requiring developers to pad ultra-short inputs.[19] Furthermore, while the model supports massive context, the KV cache requirements at 262k tokens can exceed 32 GB of memory even if the model weights are quantized, necessitating the use of context-folding strategies or reduced windows for standard 8 GB to 16 GB GPU configurations.[20]

Primary use cases and performance benchmarks

The 4B model is uniquely positioned for agentic workflows where visual understanding and tool usage are required simultaneously.[2] Official docs and community tests highlight its prowess in document recognition and OCR, where it frequently outperforms cloud models many times its

size.[22]

Task Type	Key Benchmark	Qwen3.5:4B Score	Industry Peer Comparison
Document Parsing	OmniDocBench1.5	86.2 [7]	Claude 4.5 Opus (87.7)
Visual Agent Tasks	ScreenSpot Pro	54.5 [25]	Qwen3-VL-30B (48.5)
Tool Invocation	BFCL-V4	50.3 [7]	Qwen3-30B-A3B (42.4)
General Knowledge	MMLU-Pro	79.1 [7]	GPT-OSS-120B (80.8)
Math Reasoning	MathVista	85.1 [7]	Gemini-2.5 Flash-Lite (72.8)

These benchmarks suggest that the 4B model is ideal for building local document analysis tools, e-commerce product inspectors, autonomous UI agents, and multilingual customer support bots.[14] Its strength in "spatial intelligence" makes it particularly capable of recognizing elements on a screen and outputting normalized relative coordinates for robotic or software automation.[2]

The hybrid thinking guide: syntax for dynamic reasoning

Processing Iteration 3/10: Investigating API parameters and soft-switch consistency for thinking mode. Qwen3.5:4B introduces a sophisticated dual-mode operation that allows it to function either as a high-speed instructor or a deep-reasoning agent.[11] This "Hybrid Thinking Mode" is achieved through a built-in reasoning phase where the model generates a chain-of-thought (CoT) trace before arriving at a final response.[28] For infrastructure engineers, the challenge lies in the inconsistent implementation of toggles across various inference frameworks in the weeks following the March 2026 release.[31]

API-level hard switches: enable_thinking

The primary method for controlling reasoning is the enable_thinking parameter, passed through

the extra_body or chat_template_kwargs fields in an OpenAI-compatible API call.[28] In official Alibaba Cloud Model Studio deployments, enable_thinking is a top-level parameter in the body, whereas in local servers like vLLM or SGLang, it is frequently nested within chat_template_kwargs to ensure the Jinja2 template correctly handles the <think> token insertion.[7]

When enable_thinking=True, the model operates in its reasoning-enhanced state, generating tokens inside <think>...</think> tags.[19] This mode is recommended for complex logic, math, and precise coding, but it introduces higher latency as the model may generate several hundred reasoning tokens before the actual answer begins.[34] To manage costs and latency, developers should utilize the thinking_budget parameter to cap the reasoning phase tokens.[34]

Prompt-level soft switches and community workarounds

A point of confusion in the developer community stems from the /think and /no_think soft switches.[31] While the Qwen3 series supported these as prompt-level instructions to toggle reasoning turn-by-turn, the official documentation for Qwen3.5:4B indicates that these soft switches are no longer officially supported as part of the core instruct tuning.[31] Instead, the model is designed to think by default, and disabling this behavior must be handled at the API or template level.[29]

Community benchmarkers have noted that adding /no_think to a prompt often results in the model still generating an empty <think></think> block, which can cause context-shifting issues and unnecessary reprocessing of the KV cache.[29] For a truly direct response, the most effective method is to modify the Jinja2 chat template at the server start to set enable_thinking = false globally, or to pass a custom template that removes the thinking placeholders entirely.[31]

Structured reasoning output parsing

For applications requiring programmatic access to the reasoning chain, recent versions of vLLM and SGLang have introduced reasoning parsers (--reasoning-parser qwen3).[28] These parsers extract the content between the <think> tags and return it in a dedicated reasoning_content field in the API response, separate from the content field.[28] This is critical for building transparency into AI agents, allowing users to see the "why" behind an agent's action without polluting the final answer with deliberation prose.[28]

Implementation Method	Command / Syntax	Effectiveness
API Hard Switch	"enable_thinking": false [31]	High (Native Support)

vLLM Launcher	--reasoning-parser qwen3 [28]	High (Structured Data)
Prompt Switch	/no_think [30]	Low (Broken/Inconsistent)
Jinja Template	{%- set enable_thinking = false %} [31]	Maximum (Prevents any CoT)

Infrastructure engineers are warned that "overthinking" is a persistent issue in the 4B variant, where the model may enter an existential loop when prompted with simple greetings like "Hi".[38] In multi-step RAG pipelines, the cumulative latency of unnecessary thinking can render the system impractical, making it essential to disable thinking for tasks such as query rewriting or basic routing.[38]

The few-shot matrix: when to use examples vs. zero-shot

Processing Iteration 4/10: Benchmarking few-shot resilience across task domains. Small Language Models are notoriously sensitive to prompt complexity, often exhibiting a "logic collapse" when overwhelmed by too many examples or contradictory formatting instructions.[5] Qwen3.5:4B, however, demonstrates a unique resilience compared to its 0.8B and 2B counterparts.[46] A systematic evaluation using TF-IDF example selection across classification, code fixing, and summarization tasks reveals a complex performance curve.[46]

Stable vs. collapsing tasks

The 4B model acts as the "sweet spot" for developers because it maintains consistent performance across tasks where smaller variants crater.[46] Community benchmarks show that the 0.8B model's code fixing accuracy drops from 67% to 33% the moment a single example is added, and it never recovers.[46] In contrast, the 4B variant is described as "rock solid," showing either stability or marginal gains with added shots.[46]

- **Classification Resilience**: For standard document tagging or sentiment analysis, the 4B variant is often perfect at zero-shot.[46] Adding examples in this domain follows a clean learning curve for the 0.8B model, but for the 4B variant, it primarily serves to stabilize high-cardinality label sets rather than to increase raw accuracy.[46]
- **Coding Constraints**: Developers have observed that while the 4B model handles coding examples better than the 2B model (which peaks at 1-2 shots and then falls back), it still performs best in a "Zero-Shot Thinking" configuration.[46] Providing multiple code examples can sometimes cause the model to over-fit to the style of the examples at the expense of solving the actual logical problem.[46]

Optimal shot counts for the 4B variant

Based on community and official data, the following matrix should guide prompt construction for the Qwen3.5:4B model.[46]

Task Category	Recommendation	Optimal Shot Count	Reason
Logic/Math	Zero-Shot Thinking	0-Shot [46]	Examples bypass CoT, reducing accuracy.[46]
Structure Extraction	Few-Shot (Format)	2 to 4-Shot [47]	Ensures consistent JSON/XML keys.[39]
Translation	Zero-Shot	0-Shot [46]	4B natively supports 201 languages.[6]
Code Fixing	One-Shot	1-Shot [46]	Grounds the model in the specific bug style.[46]
Summarization	Zero-Shot	0-Shot [46]	F1 score scales with model size, not shots.[46]
Tool Calling	One-Shot System	1-Shot [43]	Defines the ReAct pattern for the agent.[50]

The critical takeaway for infrastructure specialists is that "production coding needs Claude-level context," but for edge-based preprocessing and metadata extraction, the 4B model's zero-shot performance is sufficient.[46] Blindly adding shots to an SLM prompt can be counterproductive, wasting valuable context window tokens and increasing the risk of the model entering an "infinite logic loop" during the thinking phase.[46]

Multimodal syntax: unified vision-language integration

Processing Iteration 5/10: Mapping early-fusion architecture to multimodal prompting standards. Qwen3.5:4B represents a significant departure from previous vision-language

models by utilizing an "early-fusion" architecture.[6] Traditional "mid-fusion" or "late-fusion" models use a pretrained vision encoder to project images into the language model's space, which often results in a loss of fine-grained spatial data.[11] In contrast, Qwen3.5 processes image patches and text tokens in a single unified transformer stream from the earliest layers.[2]

Unified token structure and resolution

The early-fusion foundation means the model perceives visuals with the same granularity it does text.[2] It supports dynamic resolution, mapping images to a sequence of tokens ranging from 4 to 16,384 tokens.[7] This enables what Alibaba calls "pixel-perfect element detection," essential for agents operating PC or mobile GUIs.[2] In practical terms, this allows the 4B variant to succeed at document recognition tasks (OmniDocBench 90.8%) and spatial intelligence tests that previously required much larger models.[22]

Multimodal API syntax: copy-pasteable implementation

Developers integrating Qwen3.5:4B vision capabilities via OpenAI-compatible APIs must follow a specific JSON structure for message content.[7] The content field must be an array of objects specifying both the image_url and the text query.[7]

JSON

```
{
 "model": "Qwen/Qwen3.5-4B",
 "messages":
   }
 ],
 "temperature": 0.6,
 "extra_body": {
   "enable_thinking": true,
   "thinking_budget": 1500
 }
}
```

For video processing, the model handles frame extraction internally or via processors such as decord or torchcodec.[27] Official best practices for hour-scale videos suggest overriding the longest_edge parameter to enable a higher video token count (~224k tokens), which is achievable thanks to the model's DeltaNet context efficiency.[7]

Spatial grounding and bounding boxes

A unique feature of the Qwen3.5 architecture is its bbox (bounding box) output format.[26] When performing object detection or UI element recognition, the model outputs coordinates in normalized relative scale of 1000.[27] This is more robust for cross-platform automation than pixel-based coordinates, as it accounts for varying screen resolutions and aspect ratios.[2]

Grounding Task	Output Format	Accuracy Note
Object Detection	[ymin, xmin, ymax, xmax]	Scale: 0 to 1000.[27]
UI element detection	<button>...</button>	Pixel-perfect element alignment.[24]
Document Layout	Markdown with Table tags	Preserves original layout integrity.[15]

Developers using frameworks like ms-swift can provide grounding data in absolute coordinates, which the engine automatically converts to the model's native 1000-scale relative format during fine-tuning or inference.[27]

Advanced prompt templates: Copy-pasteable system messages

Processing Iteration 6/10: Compiling high-performing system prompts for agentic and specialized tasks. Because Qwen3.5:4B is optimized for "agentic" use cases, the system prompt must do more than define a persona; it must provide a structured operational framework.[9] The following templates are derived from official training documentation and March 2026 community best practices.[43]

Template 1: Autonomous GUI Agent (Multimodal)

This prompt maximizes the model's spatial reasoning and tool-calling abilities for UI navigation.[2]

Role: You are a high-fidelity Autonomous GUI Agent.

Objective: Analyze screenshots and execute precise actions to complete user requests.

Capabilities:

- Screen element identification (buttons, inputs, sliders).
- Tool use for clicks, keystrokes, and scrolling.
- Error recovery and task planning.

Operational Rules:

1. Always begin by analyzing the layout inside tags.
2. Identify interactive elements using normalized relative coordinates [0-1000].
3. Format actions as: <action=click>{{"point": }}</action>
4. If the screen has not changed after an action, assume the action failed and try an alternative approach.
5. Provide a one-sentence summary of the current progress after each turn.

Template 2: Precise Multi-Language Coder

Optimized for logic consistency and bug fixing while preventing conversational bloat.[17]

Role: Senior AI Software Engineer.

Style: Terse, technical, and precise.

Constraint:

- Use Thinking Mode for logical derivation.
- No conversational filler. Provide only the code or the fix.
- For all math or logic proofs, use \boxed{} for the final answer.

Expertise: Python (System Level), C++, Rust, and Next.js.

Instructions:

1. When reviewing code, categorize issues into: logic error, syntax error, performance bottleneck, or security risk.
2. If a bug is found, provide a minimal reproducible example (MRE) of the fix.
3. Always verify cross-platform compatibility for system-level calls.

Template 3: Structured Document Classifier (RAG Preprocessor)

Leverages the model's 90%+ accuracy in document understanding for automated pipeline classification.[12]

Role: Document Intelligence Specialist.

Objective: Classify incoming multi-format documents into the provided JSON schema.

Categories: [invoice, passport, contract, medical_report, insurance_claim, unknown]

Guidelines:

1. Analyze visual cues (logos, table structures, signatures) to determine the category.
2. Extract the confidence score (0.0 to 1.0) based on visual legibility.
3. If multiple categories apply, select the most specific one.

Output Format:

```
{
"category": "string",
"confidence": "float",
```

```
"key_fields_detected": ["field1", "field2"],
"visual_quality": "low/medium/high"
}
```

These templates are designed to be used in conjunction with the recommended hyperparameters for their respective modes, ensuring the 4B model remains within its "intelligence density" sweet spot.[38]

The hyperparameter cheat sheet: preventing loops and hallucinations

Processing Iteration 7/10: Consolidating community and official settings for stability. The greatest technical hurdle for developers running Qwen3.5:4B locally is the "Infinite Chain of Thought" repetition loop.[28] This occurs when the model gets caught in a cycle of self-correction inside the <think> block.[31] Preventing this requires an exact balance of sampling parameters and hardware-level fixes.[28]

Official recommended settings

Alibaba and the Unsloth team have published specific presets for the 4B variant.[7]

Scenario	Mode	Temperature	Top_P	Presence Penalty	Note
General Task	Thinking	1.0 [56]	0.95 [28]	1.5 [56]	Balanced creativity/logic.
Precise Coding	Thinking	0.6 [28]	0.95 [7]	0.0 [56]	Tightens logical focus.
Edge Chat	Instruct	0.7 [28]	0.80 [28]	1.5 [28]	Direct response, no CoT.
Complex RAG	Instruct	1.0 [58]	0.95 [62]	2.0 [62]	Forces variety in answers.

The key_gdiff bug and repetition mitigation

A significant community discovery in early February 2026 was a bug in llama.cpp regarding the vectorized key_gdiff calculation.[61] This bug caused MoE-based models like Qwen3.5 to miscalculate token probabilities, leading to "stumbling" outputs where the model misidentifies errors or repeats phrases.[61]

Infrastructure engineers must:

1. **Update Inference Engines**: Ensure llama.cpp or llama-server is built on version 7925 or later.[61]
2. **Re-download GGUFs**: The imatrix calibration dataset was updated on March 5, 2026, to fix these outliers.[17]
3. **Manage Temperature**: Paradoxically, temperature=0 (greedy decoding) makes loops worse in reasoning models because it locks the model into the most likely "hedge" clause.[38] A temperature of 0.6 is the official "safe" floor.[28]
4. **Context Buffering**: Infinite loops can also be triggered if the context window is set too low (e.g., 2,048 tokens). Developers are advised to maintain at least 128,000 tokens for reasoning tasks to preserve the model's internal attention state.[17]

Workflow integration and agentic orchestration

Processing Iteration 8/10: Integrating 4B agents with Model Context Protocol (MCP). The ultimate utility of Qwen3.5:4B lies in its ability to act as a local agent coordinator.[2] To facilitate this, integration with the Model Context Protocol (MCP) and orchestration frameworks like Qwen-Agent is essential.[50]

Registering MCP servers

MCP allows the model to interact with standardized tools provided by external servers (e.g., Google Search, Bright Data, local file systems).[21] The Qwen-Agent framework encapsulates the tool-calling templates required for the 4B variant.[51]

Critical Configuration Warning: March 2026 community testing has identified that registering more than two MCP servers simultaneously can break the tool-call parser for the first server.[44] This is due to prompt-bloat exceeding the model's "instruction density" limit.[44] It is recommended to use a single, consolidated MCP gateway or to rotate servers dynamically based on the user's intent.[44]

Integration with vLLM and SGLang

For production deployments, vLLM remains the gold standard for serving Qwen3.5:4B.[9] Its PagedAttention and continuous batching features are critical for handling the large KV caches

required for the model's 262k context window.[12]

Engine	Launch Parameter	Benefit
vLLM	--reasoning-parser qwen3 [36]	Correctly routes CoT to separate API fields.
SGLang	--speculative-algo NEXTN [67]	Uses Multi-Token Prediction for 3x speedup.
vLLM	--mm-processor-cache-type shm [54]	Reduces vision latency via shared memory.
llama.cpp	--rope-scaling yarn --rope-scale 4 [28]	Extends context from 32k to 128k+ natively.

Infrastructure specialists should also note that FP8 quantization is the recommended standard for serving Qwen3.5 models on NVIDIA Ada Lovelace or Hopper GPUs, offering the best balance of memory usage and serving throughput.[12]

Deployment insights and future outlook

Processing Iteration 9/10: Reviewing safety, refusals, and community reception. The March 2026 release of Qwen3.5:4B occurred amidst high-profile personnel departures from Alibaba's Qwen team, including lead Junyang Lin.[68] Despite this internal turmoil, the model series has dominated open-source rankings due to its "lossless" architecture and high performance-per-watt.[2]

Safety and uncensored variants

The official Qwen3.5:4B model includes robust guardrails, having been trained with "cybersecurity awareness" to identify vulnerabilities without facilitating exploits.[71] For research and unrestrained local testing, community members have released "Aggressive" uncensored variants (e.g., HauhauCS Aggressive).[8] These variants maintain zero refused prompts during testing (0/465) and are reported to perform better in roleplay and creative writing tasks where official guardrails sometimes cause logic loops or refusals.[8]

Intelligence density and edge AI economics

The philosophical core of the Qwen3.5 project is "More Intelligence, Less Compute".[2] By running a model with GPT-5 class multimodal understanding locally, bootstrapped startups and

individual developers can offer "unlimited usage without worrying about cost scaling".[14] This shift destroys traditional cloud AI economics, where API bills for document processing can reach hundreds of dollars per month.[14]

Feature	Qwen3.5:4B Local	Traditional Cloud API
Marginal Cost	$0 per token [14]	$0.16 to $1.30 per million [71]
Privacy	Data never leaves hardware [14]	Cloud provider logging.
Offline Capability	Full support [14]	No support.
Latency	Local network speed (sub-10ms) [14]	500ms to 2000ms+ network lag.

The 4B variant represents the ideal "on-device" foundational model, large enough to be stable but small enough to run on current-generation consumer hardware.[11]

Nuanced conclusions on SLM optimization

Processing Iteration 10/10: Final consolidation and validation of integration recommendations. Deploying Qwen3.5:4B successfully in early 2026 requires moving beyond the "chat model" mindset.[11] Developers must view it as an "AI Worker" that needs precise environmental configuration—specifically regarding the reasoning parser and the 3:1 DeltaNet/Softmax architecture.[2]

Final recommendations for optimal integration:

1. **Prioritize Thinking Mode for Precision**: For tasks where logic matters (coding, math), never disable reasoning. Use the 0.6 temperature floor to ensure stability.[28]
2. **Use Zero-Shot for Logic, Few-Shot for Format**: The "Few-Shot Paradox" is real; 4B is smart enough to reason logic zero-shot but needs examples to conform to complex output schemas.[46]
3. **Strict Coordinate Handling**: Use normalized relative coordinates [0-1000] for all visual agent tasks to ensure cross-device consistency.[26]
4. **Hardware Calibration**: Ensure the KV cache is allocated correctly for the desired context length. If OOM issues occur, reduce the window while maintaining a 128k floor for reasoning integrity.[6]

As the ecosystem matures, the 4B variant is likely to remain the dominant choice for

edge-based multimodal agents, providing a balance of reasoning depth and local performance that currently has no rival in the open-source landscape.

Works cited

1. Alibaba Unveils New AI Model Series Qwen 3.5 - Binance, accessed March 7, 2026, https://www.binance.com/en/square/post/03-03-2026-alibaba-unveils-new-ai-model-series-qwen-3-5-297427083420257
2. Alibaba Releases Open-Source Qwen3.5 Small Models for Edge Devices - MLQ.ai, accessed March 7, 2026, https://mlq.ai/news/alibaba-releases-open-source-qwen35-small-models-for-edge-devices/
3. Elon Musk praises Alibaba's Qwen model for its "impressive intelligence density" | KuCoin, accessed March 7, 2026, https://www.kucoin.com/news/flash/elon-musk-praises-alibaba-s-qwen-model-for-impressive-intelligence-density
4. Elon Musk Highlights Intelligence Density in Alibaba's Latest Qwen3.5 Release - MLQ.ai, accessed March 7, 2026, https://mlq.ai/news/elon-musk-highlights-intelligence-density-in-alibabas-latest-qwen35-release/
5. The Compression Problem: How Qwen 3.5 Small Rewrites the Rules of Edge AI - Medium, accessed March 7, 2026, https://medium.com/@AdithyaGiridharan/the-compression-problem-how-qwen-3-5-small-rewrites-the-rules-of-edge-ai-3062e8a9d5e8
6. Qwen/Qwen3.5-4B-Base - Hugging Face, accessed March 7, 2026, https://huggingface.co/Qwen/Qwen3.5-4B-Base
7. Qwen/Qwen3.5-4B - Hugging Face, accessed March 7, 2026, https://huggingface.co/Qwen/Qwen3.5-4B
8. Qwen3.5-4B Uncensored Aggressive Release (GGUF) : r/LocalLLaMA - Reddit, accessed March 7, 2026, https://www.reddit.com/r/LocalLLaMA/comments/1rjp08s/qwen354b_uncensored_aggressive_release_gguf/
9. qwen3.5-397b-a17b Model by Qwen | NVIDIA NIM, accessed March 7, 2026, https://build.nvidia.com/qwen/qwen3.5-397b-a17b/modelcard
10. unsloth/Qwen3.5-4B-GGUF · Hugging Face : r/LocalLLaMA - Reddit, accessed March 7, 2026, https://www.reddit.com/r/LocalLLaMA/comments/1rirts9/unslothqwen354bgguf_hugging_face/
11. Qwen 3.5 Small Model Series released | by Mehul Gupta | Data Science in Your Pocket, accessed March 7, 2026, https://medium.com/data-science-in-your-pocket/qwen-3-5-small-model-series-released-7a5ed34fcbb3
12. Qwen 3.5 Explained: Architecture, Upgrades Over Qwen 3, Benchmarks, and Real-World Use Cases - Medium, accessed March 7, 2026,

https://medium.com/data-science-in-your-pocket/qwen-3-5-explained-architecture-upgrades-over-qwen-3-benchmarks-and-real-world-use-cases-af38b01e9888
13. Qwen 3.5 Free API for everyone - Medium, accessed March 7, 2026, https://medium.com/data-science-in-your-pocket/qwen-3-5-free-api-for-everyone-afcf8ed3adb7
14. Qwen 3.5 Small Model Series Release: Local, private, free | STARTUP EDITION, accessed March 7, 2026, https://blog.mean.ceo/qwen-3-5-small-model-series-release/
15. Qwen3.5-9B Multimodal Guide: Running Free Image & Video AI In-House | Oflight Inc., accessed March 7, 2026, https://www.oflight.co.jp/en/columns/qwen35-9b-multimodal-business-guide
16. New Move by Qwen! Officially Open-Sources the Qwen3.5 Compact Model Series, Musk Expresses Views, accessed March 7, 2026, https://news.futunn.com/en/post/69502828/new-move-by-qwen-officially-open-sources-the-qwen3-5
17. Qwen3.5 - How to Run Locally Guide | Unsloth Documentation, accessed March 7, 2026, https://unsloth.ai/docs/models/qwen3.5
18. vLLM + Qwen-3-VL-30B-A3B is so fast : r/LocalLLaMA - Reddit, accessed March 7, 2026, https://www.reddit.com/r/LocalLLaMA/comments/1nyd512/vllm_qwen3vl30ba3b_is_so_fast/
19. Residual Thinking Embeddings: Diffusion Conditioning from Small LLMs - Hugging Face, accessed March 7, 2026, https://huggingface.co/blog/AbstractPhil/qwen-35-diffusion-potential-testing
20. qwen / qwen3.5-397b-a17b - NVIDIA API Documentation, accessed March 7, 2026, https://docs.api.nvidia.com/nim/reference/qwen-qwen3-5-397b-a17b
21. qwen3.5:4b - Ollama, accessed March 7, 2026, https://ollama.com/library/qwen3.5:4b
22. Qwen3.5: Features, Access, and Benchmarks - DataCamp, accessed March 7, 2026, https://www.datacamp.com/blog/qwen3-5
23. Qwen: Qwen3.5: Towards Native Multimodal Agents, accessed March 7, 2026, https://qwen.ai/blog?id=qwen3.5
24. Qwen3.5-397B-A17B API - Together AI, accessed March 7, 2026, https://www.together.ai/models/qwen3-5-397b-a17b
25. Qwen/Qwen3.5-2B - Hugging Face, accessed March 7, 2026, https://huggingface.co/Qwen/Qwen3.5-2B
26. Qwen/Qwen3-VL-4B-Instruct - Hugging Face, accessed March 7, 2026, https://huggingface.co/Qwen/Qwen3-VL-4B-Instruct
27. Qwen3.5 Best Practices — swift 4.1.0.dev0 documentation, accessed March 7, 2026, https://swift.readthedocs.io/en/latest/BestPractices/Qwen3_5-Best-Practice.html
28. Qwen/Qwen3-4B - Hugging Face, accessed March 7, 2026, https://huggingface.co/Qwen/Qwen3-4B
29. Qwen3 Highlights Qwen3 is the latest generation of large language models in

Qwen series, offering a comprehensive suite of dense - NVidia, accessed March 7, 2026, https://developer.nvidia.com/downloads/assets/ace/model_card/qwen3-8b-instruct.pdf
30. Qwen/Qwen3-14B - Hugging Face, accessed March 7, 2026, https://huggingface.co/Qwen/Qwen3-14B
31. Qwen 3.5 Model: Switching Between Thinking and Non-Thinking #1559 - GitHub, accessed March 7, 2026, https://github.com/lmstudio-ai/lmstudio-bug-tracker/issues/1559
32. Major update coming soon! I'm here, sorry for the delay. : r/Oobabooga - Reddit, accessed March 7, 2026, https://www.reddit.com/r/Oobabooga/comments/1rkodoh/major_update_coming_soon_im_here_sorry_for_the/
33. Using "/no_think" with HYBRID models does not work anymore · Issue #12575 - GitHub, accessed March 7, 2026, https://github.com/ollama/ollama/issues/12575
34. Qwen3-Max-Thinking: Hands-On With the Largest LLM in the World | DataCamp, accessed March 7, 2026, https://www.datacamp.com/de/tutorial/qwen3-max-thinking
35. Alibaba Cloud Model Studio:Deep thinking, accessed March 7, 2026, https://www.alibabacloud.com/help/en/model-studio/deep-thinking
36. Deployment example for a qwen3 model with hybrid thinking - vLLM Forums, accessed March 7, 2026, https://discuss.vllm.ai/t/deployment-example-for-a-qwen3-model-with-hybrid-thinking/1462
37. Setup and Fine-Tune Qwen 3 with Ollama - Codecademy, accessed March 7, 2026, https://www.codecademy.com/article/qwen-3-ollama-setup-and-fine-tuning
38. Qwen3.5 4B: overthinking to say hello. : r/LocalLLaMA - Reddit, accessed March 7, 2026, https://www.reddit.com/r/LocalLLaMA/comments/1rj8x1q/qwen35_4b_overthinking_to_say_hello/
39. How to Access the Qwen3-Max-Thinking API ? - Apidog, accessed March 7, 2026, https://apidog.com/blog/qwen3-max-thinking-api/
40. Final Qwen3.5 Unsloth GGUF Update! : r/LocalLLaMA - Reddit, accessed March 7, 2026, https://www.reddit.com/r/LocalLLaMA/comments/1rlkptk/final_qwen35_unsloth_gguf_update/
41. vLLM - Qwen, accessed March 7, 2026, https://qwen.readthedocs.io/en/latest/deployment/vllm.html
42. [Bug]: `qwen3` reasoning parser incorrectly parse reasoning-only output as `content` · Issue #35221 · vllm-project/vllm - GitHub, accessed March 7, 2026, https://github.com/vllm-project/vllm/issues/35221
43. Qwen3 Instruct vs Thinking vs Coder: Model Selection Guide - Fireworks AI, accessed March 7, 2026, https://fireworks.ai/blog/qwen-3-decoded
44. PSA: LM Studio's parser silently breaks Qwen3.5 tool calling and reasoning: a year

of connected bug reports - Reddit, accessed March 7, 2026, https://www.reddit.com/r/LocalLLaMA/comments/1riwhcf/psa_lm_studios_parser_silently_breaks_qwen35_tool/
45. Qwen3.5 122B in 72GB VRAM (3x3090) is the best model available at this time — also it nails the “car wash test” - Reddit, accessed March 7, 2026, https://www.reddit.com/r/LocalLLaMA/comments/1rf2ulo/qwen35_122b_in_72gb_vram_3x3090_is_the_best_model/
46. Benchmarked Qwen 3.5 small models (0.8B/2B/4B/9B) on few-shot learning — adding examples to 0.8B code tasks actually makes it worse : r/LocalLLaMA - Reddit, accessed March 7, 2026, https://www.reddit.com/r/LocalLLaMA/comments/1rjbw0p/benchmarked_qwen_35_small_models_08b2b4b9b_on/
47. Benchmarking Small Language Models and Small Reasoning Language Models on System Log Severity Classification - arXiv, accessed March 7, 2026, https://arxiv.org/html/2601.07790v1
48. Qwen3.5 Fine-Tuning Guide - Hacker News, accessed March 7, 2026, https://news.ycombinator.com/item?id=47246296
49. Qwen3 vs Qwen3.5 performance : r/LocalLLaMA - Reddit, accessed March 7, 2026, https://www.reddit.com/r/LocalLLaMA/comments/1rlckan/qwen3_vs_qwen35_performance/
50. Function Calling - Qwen, accessed March 7, 2026, https://qwen.readthedocs.io/en/latest/framework/function_call.html
51. How to Integrate Qwen-Agent with MCP for Real-Time AI Agents - Bright Data, accessed March 7, 2026, https://brightdata.com/blog/ai/qwen-agent-with-bright-data-mcp-server
52. Can't use Claude Code with Ollama local model qwen3.5:35b-a3b-q4_K_M - Reddit, accessed March 7, 2026, https://www.reddit.com/r/LocalLLaMA/comments/1rhgzyb/cant_use_claude_code_with_ollama_local_model/
53. Phi-4-reasoning-vision and the lessons of training a multimodal reasoning model - Microsoft, accessed March 7, 2026, https://www.microsoft.com/en-us/research/blog/phi-4-reasoning-vision-and-the-lessons-of-training-a-multimodal-reasoning-model/
54. Qwen3.5 Usage Guide - vLLM Recipes, accessed March 7, 2026, https://docs.vllm.ai/projects/recipes/en/latest/Qwen/Qwen3.5.html
55. README.md · BennyDaBall/qwen3-4b-Z-Image-Engineer at main - Hugging Face, accessed March 7, 2026, https://huggingface.co/BennyDaBall/qwen3-4b-Z-Image-Engineer/blob/main/README.md
56. System prompt for Qwen3.5 (27B/35BA3B) to reduce overthinking? : r/LocalLLaMA - Reddit, accessed March 7, 2026, https://www.reddit.com/r/LocalLLaMA/comments/1rg0487/system_prompt_for_qwen35_27b35ba3b_to_reduce/
57. Qwen/Qwen3-4B-Instruct-2507 - Hugging Face, accessed March 7, 2026,

https://huggingface.co/Qwen/Qwen3-4B-Instruct-2507
58. Qwen 3.5 9B 4-bit tuning suggestions to avoid infinite thinking loops? : r/Qwen_AI - Reddit, accessed March 7, 2026, https://www.reddit.com/r/Qwen_AI/comments/1rjywf1/qwen_35_9b_4bit_tuning_suggestions_to_avoid/
59. Qwen3.5 is the large language model series developed by Qwen team, Alibaba Cloud. - GitHub, accessed March 7, 2026, https://github.com/QwenLM/Qwen3.5
60. README.md · Qwen/Qwen3-4B at eb971e9fb1f41c13b5e5a56e56886305c5ad94a0 - Hugging Face, accessed March 7, 2026, https://huggingface.co/Qwen/Qwen3-4B/blob/eb971e9fb1f41c13b5e5a56e56886305c5ad94a0/README.md
61. Qwen3-Coder-Next: How to Run Locally | Unsloth Documentation, accessed March 7, 2026, https://unsloth.ai/docs/models/qwen3-coder-next
62. HOW TO FIX QWEN3.5 OVERTHINK : r/LocalLLaMA - Reddit, accessed March 7, 2026, https://www.reddit.com/r/LocalLLaMA/comments/1rjsgy6/how_to_fix_qwen35_overthink/
63. Eval bug: Qwen3-Coder-Next Poor Outputs · Issue #19305 · ggml-org/llama.cpp - GitHub, accessed March 7, 2026, https://github.com/ggml-org/llama.cpp/issues/19305
64. Qwen3 - How to Run & Fine-tune | Unsloth Documentation, accessed March 7, 2026, https://unsloth.ai/docs/models/qwen3-how-to-run-and-fine-tune
65. Help - Qwen3 keeps repeating itself and won't stop : r/LocalLLaMA - Reddit, accessed March 7, 2026, https://www.reddit.com/r/LocalLLaMA/comments/1kccjd7/help_qwen3_keeps_repeating_itself_and_wont_stop/
66. The Best Open-Source LLMs in 2026 - Bento, accessed March 7, 2026, https://www.bentoml.com/blog/navigating-the-world-of-open-source-large-language-models
67. Qwen/Qwen3.5-0.8B - Hugging Face, accessed March 7, 2026, https://huggingface.co/Qwen/Qwen3.5-0.8B
68. Archive for Wednesday, 4th March 2026 - Simon Willison's Weblog, accessed March 7, 2026, https://simonwillison.net/2026/Mar/4/
69. Elon Musk Praised Alibaba's LLM. Then Its Leader Quit: 'Bye My Beloved Qwen' - Finviz, accessed March 7, 2026, https://finviz.com/news/329025/elon-musk-praised-alibabas-llm-then-its-leader-quit-bye-my-beloved-qwen
70. Alibaba has released 4 new Qwen3.5 models from 0.8B to 9B. 9B version easily runs on standard PC, and scores higher in Artificial Analysis index than ChatGPT's o1 model did. : r/singularity - Reddit, accessed March 7, 2026, https://www.reddit.com/r/singularity/comments/1rlyz3l/alibaba_has_released_4_new_qwen35_models_from_08b/
71. Models - OpenRouter, accessed March 7, 2026, https://openrouter.ai/models
72. How to Use Qwen 3.5 Small Model Series - Apidog, accessed March 7, 2026,

https://apidog.com/blog/use-qwen-3-5-small-models/